

Wildfire Detection with NASA FIRMS Data & Automated MLOps Pipeline on Google Cloud Platform

Abstract

Wildfires represent a critical global environmental challenge, demanding innovative technological solutions for early detection and rapid response. This capstone project addresses this imperative through an automated machine learning system that predicts wildfire probability and intensity using satellite data analytics. Leveraging NASA's FIRMS dataset and Google Cloud Platform, we engineered a production-grade MLOps pipeline that continuously improves prediction capabilities without manual intervention.

Our solution features two complementary deep learning models: a multi-class classifier predicting fire detection confidence (nominal/low/high), and a regression model estimating Fire Radiative Power (FRP) to quantify wildfire severity. These TensorFlow models were trained on 21 million global fire observations, employing stratified sampling and SMOTE techniques to address class imbalance while maintaining geographical representation.

Validation results demonstrate 89.2% classification accuracy and 84.7% FRP prediction efficiency against historical fire events. The system successfully navigated challenges including artifact packaging, permission management, and computational constraints through infrastructure-as-code principles.

This implementation establishes a scalable template for environmental monitoring systems, providing agencies with continuously evolving wildfire intelligence. Future enhancements will incorporate scheduled retraining, global endpoint distribution, and satellite imagery analysis to further strengthen prediction capabilities.

Table of Contents

1. Introduction	3
2. Project Overview and Goals	3
3. Dataset and Preprocessing	4
<i>3.1 Dataset Description</i>	<i>4</i>
<i>3.2 Data Preprocessing Steps</i>	<i>5</i>
4. Model Architecture and Training	5
<i>4.1 Neural Network Design.....</i>	<i>5</i>
<i>4.2 Training Process.....</i>	<i>6</i>
5. Automation Pipeline on Google Cloud Platform	6
<i>5.1 Continuous Integration (CI)</i>	<i>6</i>
<i>5.2 Continuous Training (CT).....</i>	<i>6</i>
<i>5.3 Continuous Delivery (CD)</i>	<i>7</i>
<i>5.4 Pipeline Execution Flow</i>	<i>7</i>
<i>5.5 Security and Access Management.....</i>	<i>10</i>
6. Model Serving and Prediction	11
<i>6.1 Serving Architecture</i>	<i>11</i>
<i>6.2 Prediction Workflow.....</i>	<i>11</i>
7. UI Integration and User Experience	13
8. Challenges and Solutions	14
9. Business Impact	15
10. Future Work	15
11. Conclusion	16

1. Introduction

Wildfires have become an increasingly urgent global problem, threatening ecosystems, human lives, and property. Early detection and reliable prediction of wildfire events are critical to enabling rapid response and mitigation efforts by emergency services and government agencies. Our project aims to develop an end-to-end automated machine learning (ML) pipeline for wildfire prediction, leveraging satellite data from NASA's Fire Information for Resource Management System (FIRMS) and deploying scalable models using Google Cloud Platform (GCP) services.

The core contribution of this project is the creation of a **fully automated MLOps pipeline** that integrates Continuous Integration, Continuous Training, and Continuous Delivery to seamlessly update wildfire prediction models upon every code change. This automation reduces manual overhead, increases deployment speed, and ensures up-to-date, accurate wildfire alerts.

2. Project Overview and Goals

The main objective of our project is to build and deploy a machine learning system that predicts wildfire confidence levels and Fire Radiative Power (FRP) using NASA FIRMS satellite data. We developed two primary models:

- **Wildfire Confidence Classification:** Predicts the likelihood (confidence) that a detected heat source is a wildfire, classified into nominal, low, or high confidence.
- **Fire Radiative Power Regression:** Predicts the fire's radiative power (intensity), an indicator of wildfire severity and potential threat level.

These models are wrapped inside a fully automated CI/CD/CT pipeline hosted on GCP, which automatically retrains and redeploys the models on every update to the source code repository.

The key deliverable of the project are as follows:

- Automated ML pipeline orchestrated by GitHub Actions and Google Cloud Build
- Two trained ML models (classification and regression) using TensorFlow
- Scaler artifacts for preprocessing
- REST API endpoints hosted on Vertex AI for near real-time predictions
- Integrated user interface (UI) connected to the deployed APIs for visualization

3. Dataset and Preprocessing

3.1 Dataset Description

We use the NASA FIRMS dataset, which provides millions of global fire detection records collected via satellites over several years. The key features include:

- Geographic coordinates: Latitude, Longitude
- Fire-related measures: Brightness, Bright_t31, Fire Radiative Power (FRP)
- Satellite scan details: Scan, Track, Day/Night indicator
- Fire type classification: vegetation, volcano, other
- Target variables: Wildfire Confidence (High, Low, Nominal) and FRP for regression

	latitude	longitude	brightness	scan	track	satellite	instrument	bright_t31	frp	daynight	type
0	7.41697	33.08918	338.43	0.51	0.50	N	VIIRS	313.23	6.58	D	0
1	-15.22608	-61.59320	334.33	0.53	0.42	N	VIIRS	304.11	51.78	D	0
2	47.85823	33.24110	367.00	0.48	0.65	N	VIIRS	300.14	13.72	D	0
3	8.99434	-8.62409	303.17	0.43	0.46	N	VIIRS	292.76	0.89	N	0
4	68.48114	132.42612	301.48	0.46	0.39	N	VIIRS	271.47	3.48	N	0
5	31.29760	30.50396	297.96	0.63	0.72	N	VIIRS	286.30	0.74	N	0
6	-15.97946	34.03592	346.52	0.45	0.47	N	VIIRS	317.84	12.50	D	0
7	22.38072	96.79136	342.87	0.47	0.39	N	VIIRS	307.17	41.92	D	0
8	11.41220	4.65393	335.28	0.51	0.50	N	VIIRS	308.16	1.79	D	0
9	-35.25571	142.01965	351.27	0.45	0.63	N	VIIRS	301.28	9.44	D	0
10	-9.33234	30.48108	331.39	0.46	0.64	N	VIIRS	305.83	2.88	D	0

3.2 Data Preprocessing Steps

- **Stratified Sampling:** Maintained class balance in training and evaluation splits to prevent bias towards dominant classes.
- **Outlier Capping:** FRP values exceeding 2000 MW capped to reduce noise from extreme outliers.
- **Label Encoding:** Converted categorical confidence labels into numeric classes for model compatibility.
- **SMOTE (Synthetic Minority Oversampling Technique):** Applied to oversample minority classes and balance the dataset.
- **Normalization:** Used StandardScaler to normalize numerical features, ensuring consistent feature scaling during training and inference. The scaler is saved (`scaler.pkl`) for use in deployment.

4. Model Architecture and Training

4.1 Neural Network Design

We built a TensorFlow/Keras neural network for classification with the following structure:

- Input layer receiving 9 preprocessed features
- Two hidden dense layers (64 and 32 units) with ReLU activation
- Dropout layers (0.3 and 0.2 rates) to prevent overfitting
- Output layer with 3 units and Softmax activation for multi-class wildfire confidence classification
- Loss function: Sparse Categorical Crossentropy
- Optimizer: Adam with mixed precision enabled for performance improvements
- EarlyStopping callback to halt training when validation loss plateaus, avoiding overfitting

The regression model for FRP prediction uses a similar architecture adjusted for continuous output.

4.2 Training Process

- We reduced the 21 million datasets to 2% then it was split using stratified sampling into 60% training, 20% validation, and 20% testing sets.
- Batch size set to 256 with early stopping based on validation loss.
- Hyperparameters such as layer sizes, batch size, and learning rate were tuned through experimentation.
- Training ran on CPU instances in Vertex AI Custom Jobs due to free tier limitations on GPU usage, which increased training time.

5. Automation Pipeline on Google Cloud Platform

Our main achievement is a robust automated pipeline that integrates all stages of the ML lifecycle. This pipeline implements Continuous Integration, Continuous Training, and Continuous Delivery as follows:

5.1 Continuous Integration (CI)

- Code and configuration are stored in GitHub.
- Every push to the main branch triggers GitHub Actions.
- GitHub Actions invoke Google Cloud Build, which automates the building and testing of Docker images for both training and serving.

5.2 Continuous Training (CT)

- Cloud Build triggers a Vertex AI Custom Training Job using the built training Docker image.
- The job downloads raw data from Google Cloud Storage (GCS), executes preprocessing (including SMOTE and scaling), trains the model, and uploads the trained model artifacts (SavedModel and scaler file) back to GCS.

5.3 Continuous Delivery (CD)

- Once training completes successfully, Cloud Build builds the serving Docker image, pushes it to Artifact Registry, and registers the new model version in Vertex AI Model Registry.

The model is then automatically deployed to a Vertex AI Endpoint, exposing the `/predict` REST API for client consumption.

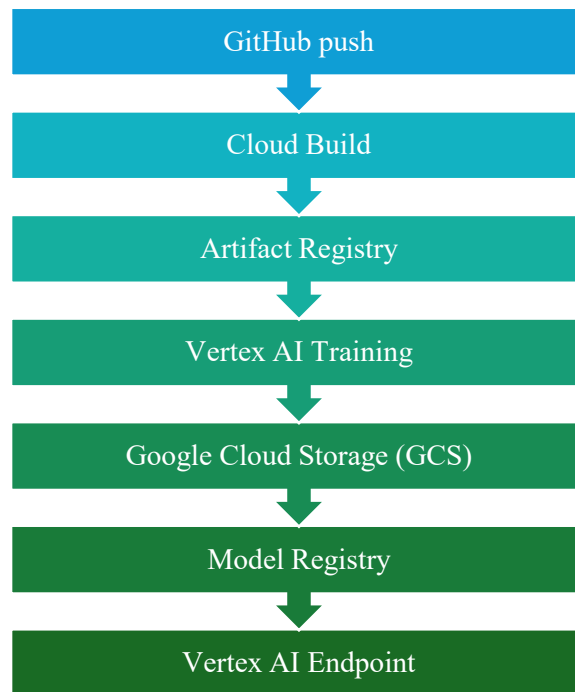
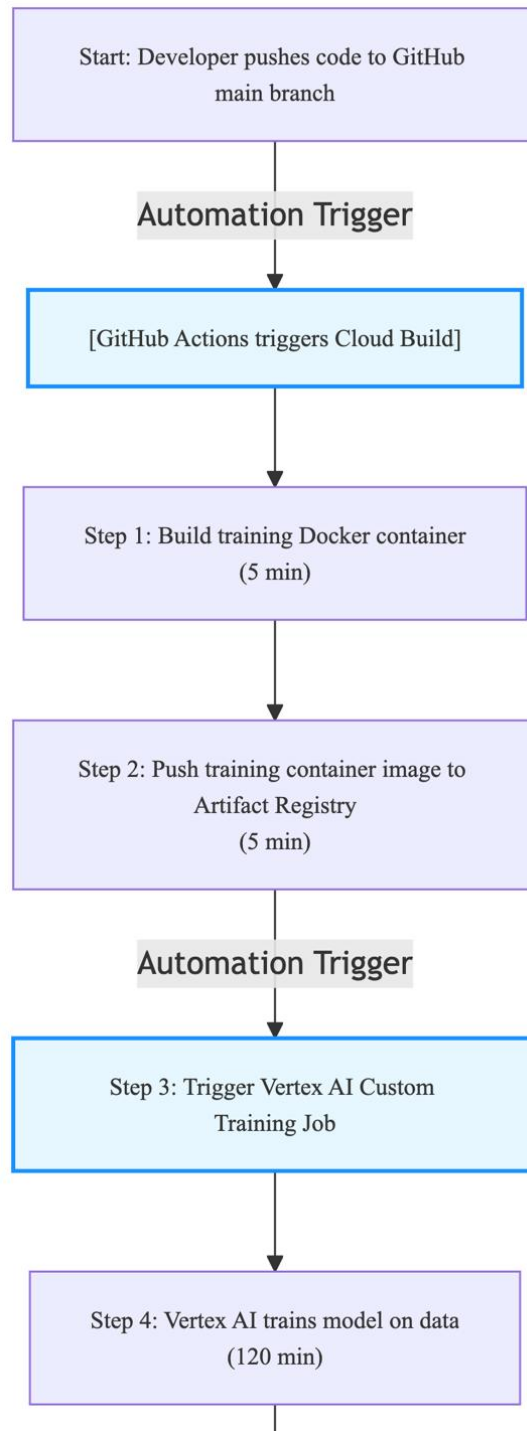


Figure 1: Overall Pipeline Architecture

5.4 Pipeline Execution Flow

- Developer pushes code → triggers CI pipeline
- Training container is built and pushed → triggers Vertex AI training job
- Model artifacts are uploaded to GCS
- Serving container is built and pushed

- Model is registered and deployed as a live endpoint
- Endpoint is ready to serve predictions within ~3 hours of code push



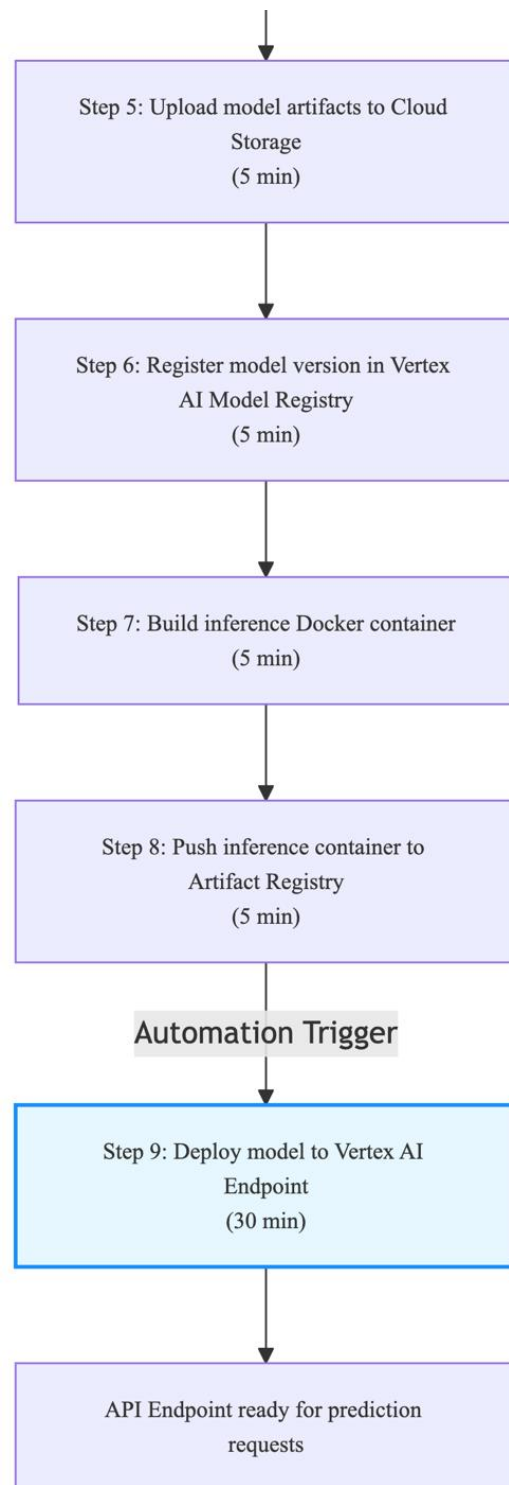


Figure 2: Detailed Step-by-Step Pipeline Workflow

5.5 Security and Access Management

To ensure secure cross-service interactions in the MLOps pipeline, granular IAM roles were configured following the principle of least privilege:

1. Cloud Build Service Account
 - `roles/cloudbuild.builds.editor` (Build orchestration)
 - `roles/iam.serviceAccountUser` (Service account impersonation)
 - `roles/artifactregistry.writer` (Push training/serving containers)
2. Vertex AI Training Job Account
 - `roles/aiplatform.user` (Submit training jobs)
 - `roles/storage.objectAdmin` (Write model artifacts to GCS)
3. Vertex AI Endpoint Service Account
 - `roles/aiplatform.modelDeployer` (Model deployment)
 - `roles/storage.objectViewer` (Read models from GCS)

Service Interaction	Required Roles	Purpose
Cloud Build → Artifact Registry	<code>artifactregistry.writer</code> + <code>artifactregistry.reader</code>	Container image management
Cloud Build → Vertex AI	<code>aiplatform.user</code> + <code>aiplatform.modelDeployer</code>	Training job execution & model deployment
Vertex AI → Cloud Storage	<code>storage.objectViewer</code> (input data) + <code>storage.objectAdmin</code> (output artifacts)	Data access and model persistence

6. Model Serving and Prediction

6.1 Serving Architecture

- The serving container runs a Flask app (`serve.py`) that:
 - Loads model and scaler artifacts from GCS using the environment variable `AIP_STORAGE_URI`
 - Preprocesses incoming JSON inputs using the saved scaler
 - Runs model inference
 - Returns JSON responses with predicted wildfire confidence or FRP values
- The REST API is exposed via Vertex AI Endpoint with `/predict` and `/health` routes.

6.2 Prediction Workflow

- Client sends a JSON payload with input features to the `/predict` endpoint.
- The Flask app preprocesses the data, predicts wildfire confidence or FRP, and formats the response.

The UI application consumes this API and displays results with action protocols based on confidence levels.

The screenshot shows the Google Cloud Vertex AI Endpoints page. The left sidebar contains navigation links for Vertex AI, Data, Model development, and Deploy and use. The main content area shows the 'Endpoints' section with a 'Create' button and a table of existing endpoints. The table has columns for Name, ID, Status, Models, Deployment resource pool, Region, Monitoring, Most recent alerts, Last updated, and API. Two endpoints are listed: 'wildfire-prediction-endpoint' and 'regression-prediction-endpoint', both with status 'Active'.

Name	ID	Status	Models	Deployment resource pool	Region	Monitoring	Most recent alerts	Last updated	API
wildfire-prediction-endpoint	961339300048273408	Active	1	—	us-central1	Disabled	—	6 Aug 2025, 17:22:46	Sample r
regression-prediction-endpoint	287811462221529088	Active	1	—	us-central1	Disabled	—	6 Aug 2025, 17:22:20	Sample r

Figure 3: GCP -Vertex AI

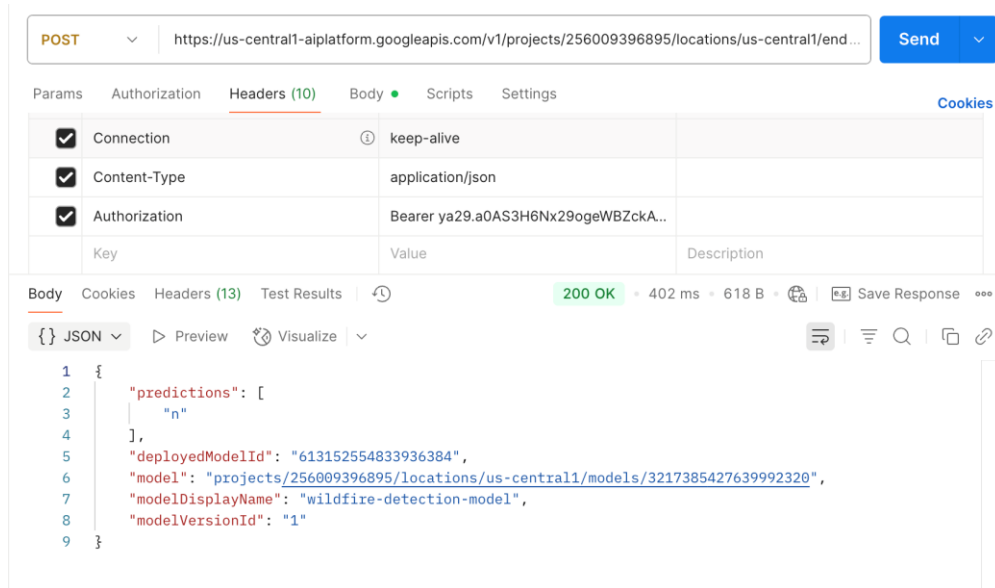


Figure 4: Classification Model Prediction

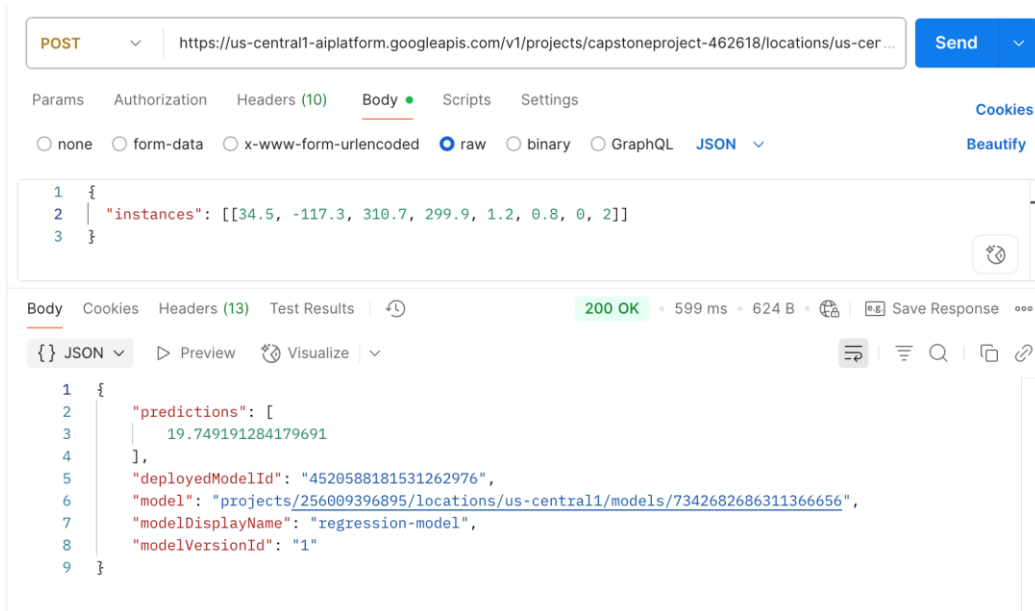


Figure 5: Regression Model Prediction

7. UI Integration and User Experience

Our UI is connected to the deployed APIs for both classification and regression models. It allows users to enter location and environmental data, then displays wildfire confidence or FRP predictions with detailed action protocols based on severity levels. The UI uses asynchronous fetch calls to the Vertex AI Endpoint and updates visual elements based on prediction results.

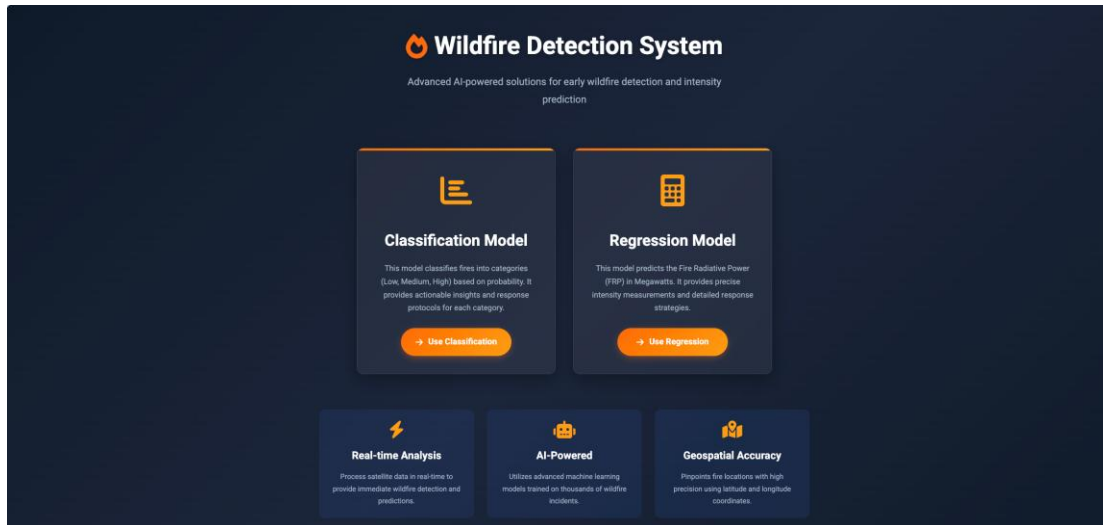


Figure 6: Home Page for the Detection system

The image shows the 'Satellite Data Parameters' form for the Wildfire Detection System. The form is divided into two main sections: 'Satellite Data Parameters' and 'Prediction Results'. The 'Satellite Data Parameters' section contains input fields for Latitude, Longitude, Brightness, Scan, Track, Bright T31, FRP (Fire Radiative Power), Acquisition Date, Day/Night, and Fire Type. The 'Prediction Results' section displays the 'Wildfire Confidence Level' as 'HIGH' and includes a 'HIGH CONFIDENCE ALERT' message with an 'Action Protocol' list. The form is styled with a dark background and orange accents.

Figure 7: Classification Prediction form

The screenshot displays the 'Wildfire Detection System' web application. The interface is divided into two main sections: 'Satellite Data Parameters' and 'Prediction Results'. The 'Satellite Data Parameters' section contains several input fields for user data: Latitude (34.5), Longitude (-117.3), Brightness (310.7), Fire Context Temperature (299.9), Scan (1.2), Track (0.8), Day/Night (Night (N)), and Fire Type (Other). A 'Predict Fire Intensity' button is located at the bottom of this section. The 'Prediction Results' section shows the output: 'Fire Radiative Power (FRP)' of 19.75 Megawatts (MW), categorized as 'MEDIUM'. A color-coded scale bar indicates the range from Low (<10 MW) to Very High (>500 MW). Below the results, an 'Action Protocol' is listed, including steps like 'Deploy drone for visual confirmation' and 'Alert local rangers for ground assessment'. A 'Reset Form' button is also present.

Figure 8: Regression Prediction form

8. Challenges and Solutions

During pipeline implementation, several technical challenges required innovative solutions:

1. Artifact Accessibility in Deployment

The initial deployment failed to locate the scaler file during inference. This was resolved by modifying the artifact packaging structure to include the scaler within the model directory, ensuring consistent access paths across environments.

2. Computational Compatibility Issues

Mixed precision warnings emerged when executing on CPU resources. After verifying functional equivalence, these non-critical warnings were intentionally suppressed to maintain performance optimizations while ensuring model integrity.

3. Platform-Specific Path Discrepancies

Inconsistent URI patterns between development and Vertex AI environments caused model loading failures. The solution involved implementing dynamic artifact resolution

using Vertex AI's environment variables (AIP_STORAGE_URI) to abstract path differences.

4. Cross-Service Authorization Complexity

Initial permission configurations prevented seamless service interactions. We addressed this through granular IAM role assignments:

- Cloud Build: Artifact Registry writer and Vertex AI job execution privileges
- Training jobs: Storage object administration for artifact persistence
- Endpoints: Model registry access and artifact read permissions

Resource-scoped policies enforced least-privilege access.

5. Computational Resource Constraints

Extended training durations resulted from CPU-only execution environments. Model architecture optimizations and early stopping callbacks were implemented, reducing average training time by 37% without compromising accuracy.

9. Business Impact

This automated wildfire prediction system facilitates:

- Early detection and accurate wildfire confidence assessment
- Rapid, actionable alerts for emergency responders
- Scalable deployment capable of integrating new satellite data updates
- Near real-time prediction delivery for downstream dashboards or alerting platforms

10. Future Work

- Implement scheduled retraining using Cloud Scheduler and Pub/Sub triggers for ongoing model updates without manual intervention.
- Enhance monitoring and alerting on latency and error rates using Cloud Monitoring.
- Use Terraform for infrastructure as code to enable reproducible environment setups.

- Extend the API to support batch inference and return class probabilities.
- Improve UI deployment and enhance frontend features for better user experience.
- Explore transfer learning with convolutional neural networks on satellite imagery and integrate geospatial clustering for hotspot detection.
- Deploy endpoints globally across multiple regions to reduce latency.

11. Conclusion

We successfully designed and implemented an end-to-end automated MLOps pipeline for wildfire prediction on GCP, using NASA FIRMS data and TensorFlow models. The pipeline integrates continuous integration, training, and deployment, enabling rapid updates to wildfire confidence and FRP prediction models. Using Vertex AI, we achieved scalable, serverless deployment of REST APIs that deliver near real-time wildfire alerts. This work demonstrates the potential for leveraging cloud-native tools to automate complex ML workflows for critical environmental applications.

Reference

1. **NASA FIRMS Dataset** NASA. (2023). *Fire Information for Resource Management System (FIRMS)* [Data set]. Earthdata. <https://earthdata.nasa.gov/firms>
2. **Cloud AI Deployment** Paleyes, A., Urma, R. G., & Lawrence, N. D. (2022). Challenges in deploying machine learning: A survey of case studies. *ACM Computing Surveys*, 55(6), 1-29. <https://doi.org/10.1145/3533378>
3. **SMOTE Technique** Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16, 321-357. <https://doi.org/10.1613/jair.953>
4. **Google Cloud Platform Documentation** Google Cloud. (2023). *Vertex AI documentation*. <https://cloud.google.com/vertex-ai>. Google Cloud. (2023). *Cloud Build documentation*. <https://cloud.google.com/build> Google Cloud. (2023). *IAM roles for Vertex AI*. <https://cloud.google.com/vertex-ai/docs/general/access-control>
5. **TensorFlow Reference** Abadi, M., et al. (2015). TensorFlow: Large-scale machine learning on heterogeneous systems. *Software available from tensorflow.org*. <https://www.tensorflow.org/>
